

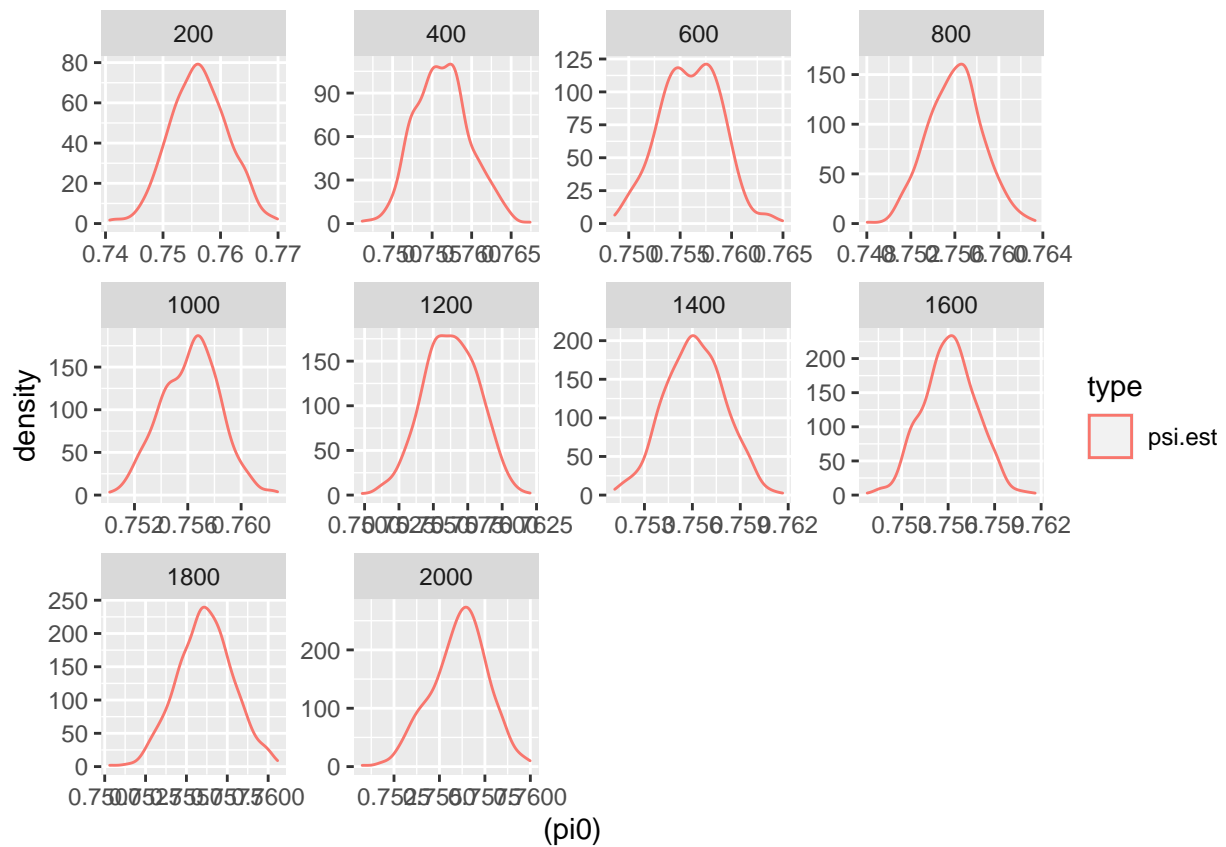
# Simulation result

## 1. case 1

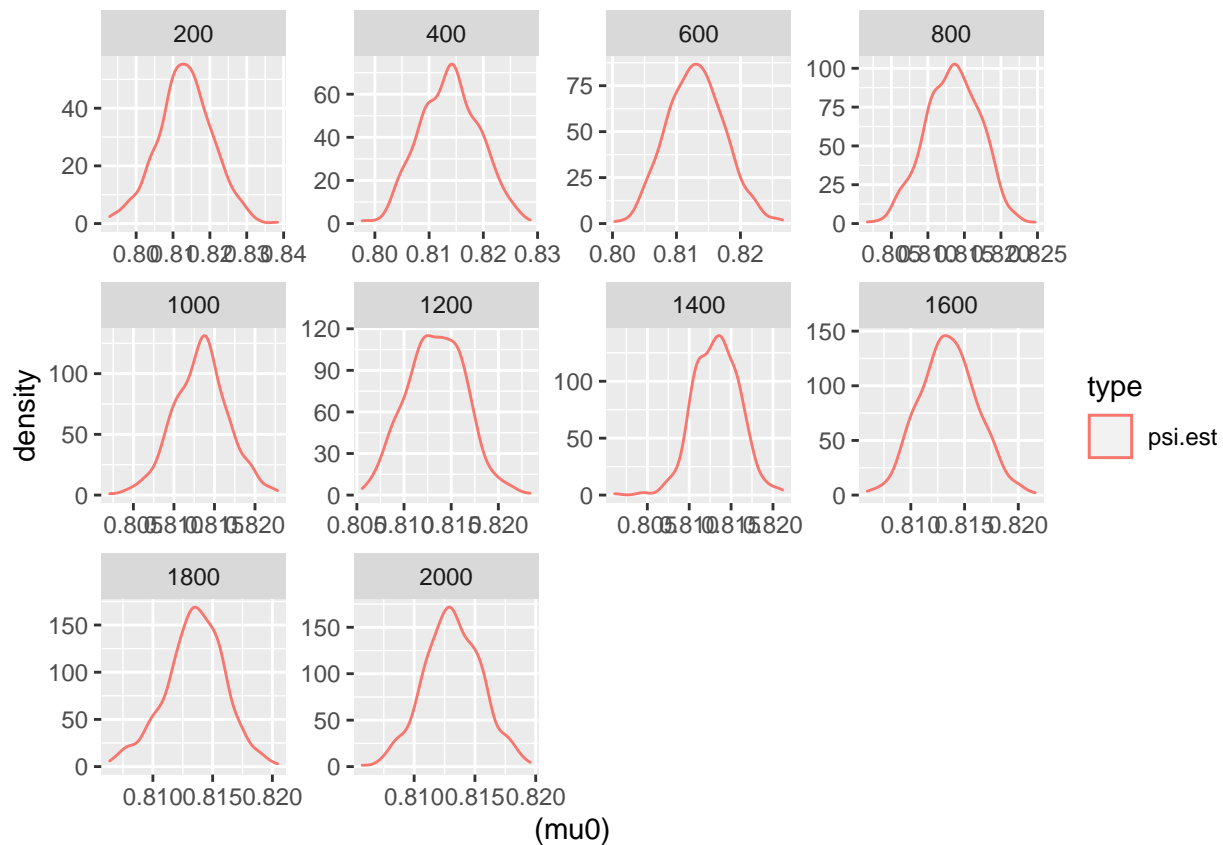
```
## true data background
# W <- matrix(runif(n*3, 0, 1), ncol=3)
# A <- rbinom(n, size = 1, prob = pi0(W))
# if(null.sims) {
#   Y <- rbinom(n, size = 1, prob = mu0.null(A, W))
#   psi0 <- mean((mu0.null(1,W) - mu0.null(0,W))^2)
#   theta0 <- var((mu0.null(1,W) - mu0.null(0,W)))
# } else {
#   Y <- rbinom(n, size = 1, prob = mu0(A, W))
#   psi0 <- mean((mu0(1,W) - mu0(0,W))^2)
#   theta0 <- var((mu0(1,W) - mu0(0,W)))
# }

# ests.sim.1 <- ests.sim(200*c(1:10), 1:500, control = list(conf.int = TRUE), null.sims=FALSE, out.glm=

load("ests.sim.1.Rdata")
ggplot(subset(ests.sim.1, (type %in% 'psi.est')) +
  geom_density(aes((pi0), color=type)) +
  facet_wrap(~n, scales='free'))
```



```
ggplot(subset(ests.sim.1, (type %in% 'psi.est'))) +
  geom_density(aes((mu0), color=type)) +
  facet_wrap(~n, scales='free')
```



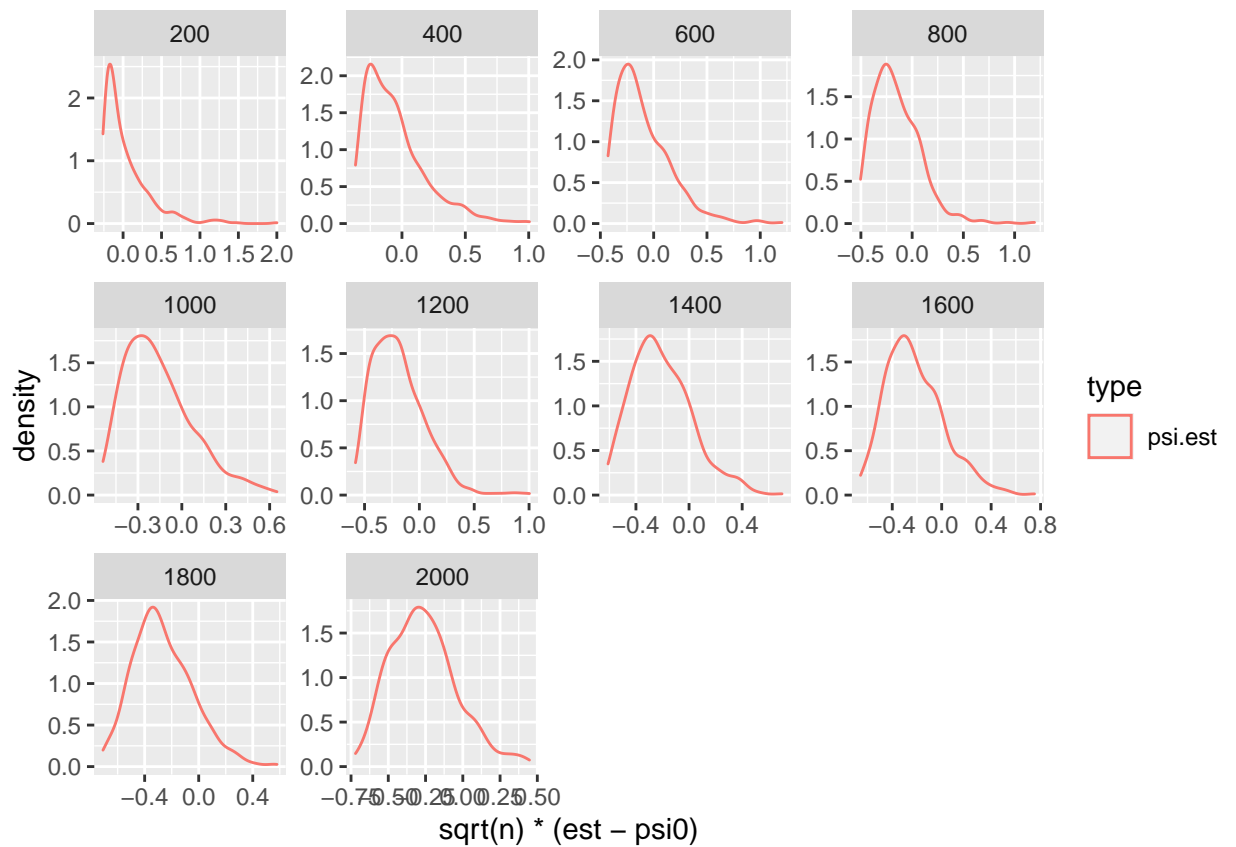
- `null.sims=FALSE`, which means that `mu1.hat` should be different from `mu0.hat`.
- The mean of `pi0` is around 0.75 and the mean of `mu0` is around 0.8.
- use `glm` to estimate both `pi.hat` and `mu.hat` in simple unit tests.

### 1.1 rates of convergence

$$n^{\frac{1}{2}}(\psi_n - \psi_0) \xrightarrow{d} N(0, \sigma_{0,\psi}^2)$$

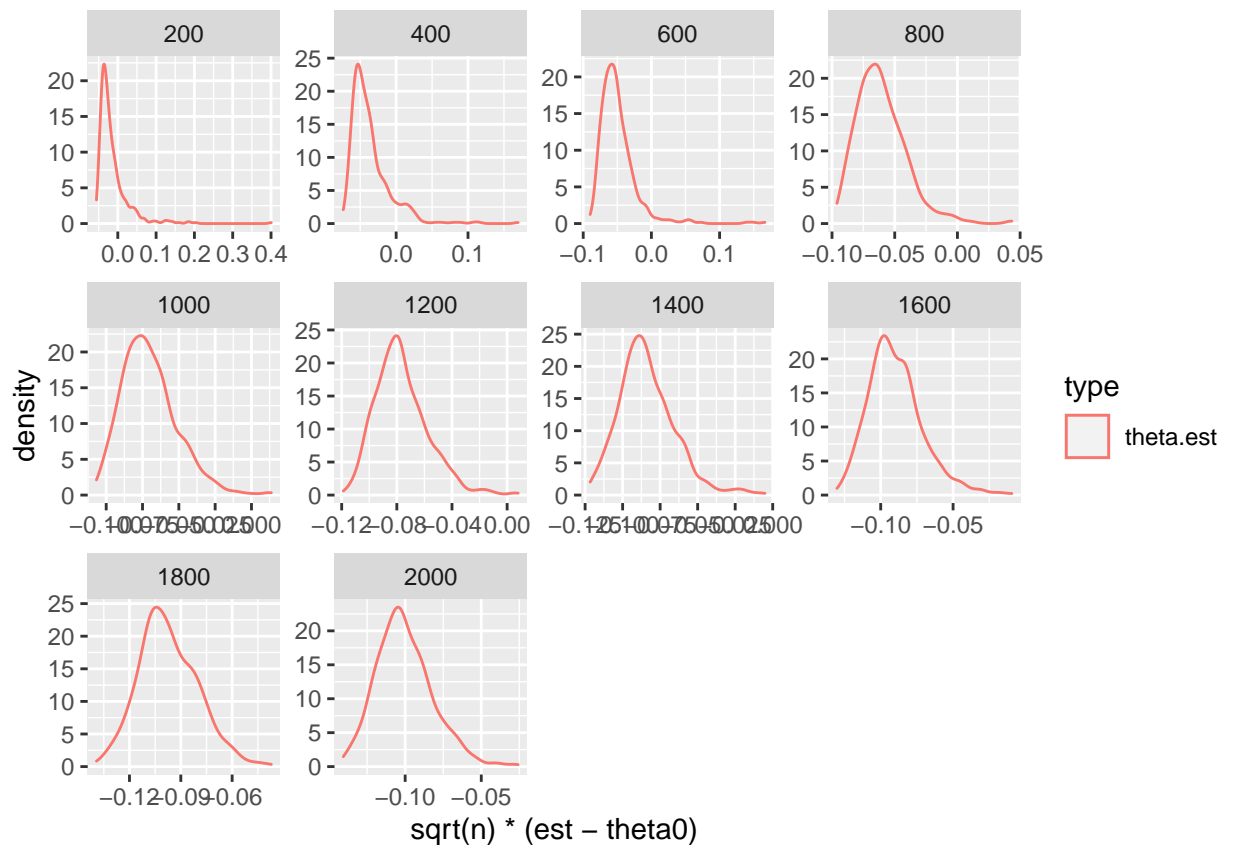
$$n^{\frac{1}{2}}(\theta_n - \theta_0) \xrightarrow{d} N(0, \sigma_{0,\theta}^2)$$

```
ggplot(subset(ests.sim.1, (type %in% 'psi.est'))) +
  geom_density(aes(sqrt(n) * (est - psi0), color=type)) +
  facet_wrap(~n, scales='free')
```



### 1.1.1.1 psi

```
ggplot(subset(ests.sim.1, (type %in% 'theta.est'))) +
  geom_density(aes(sqrt(n) * (est - theta0), color=type)) +
  facet_wrap(~n, scales='free')
```

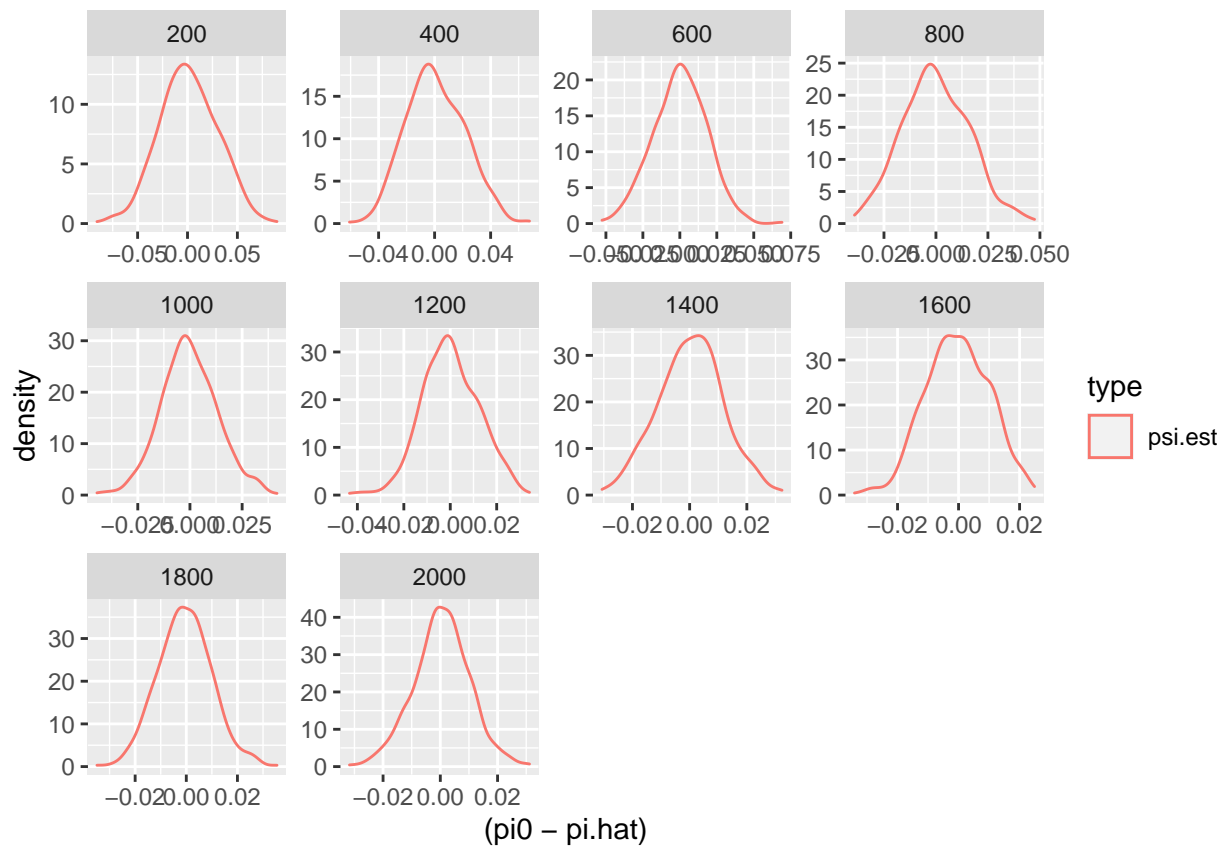


### 1.1.2 theta

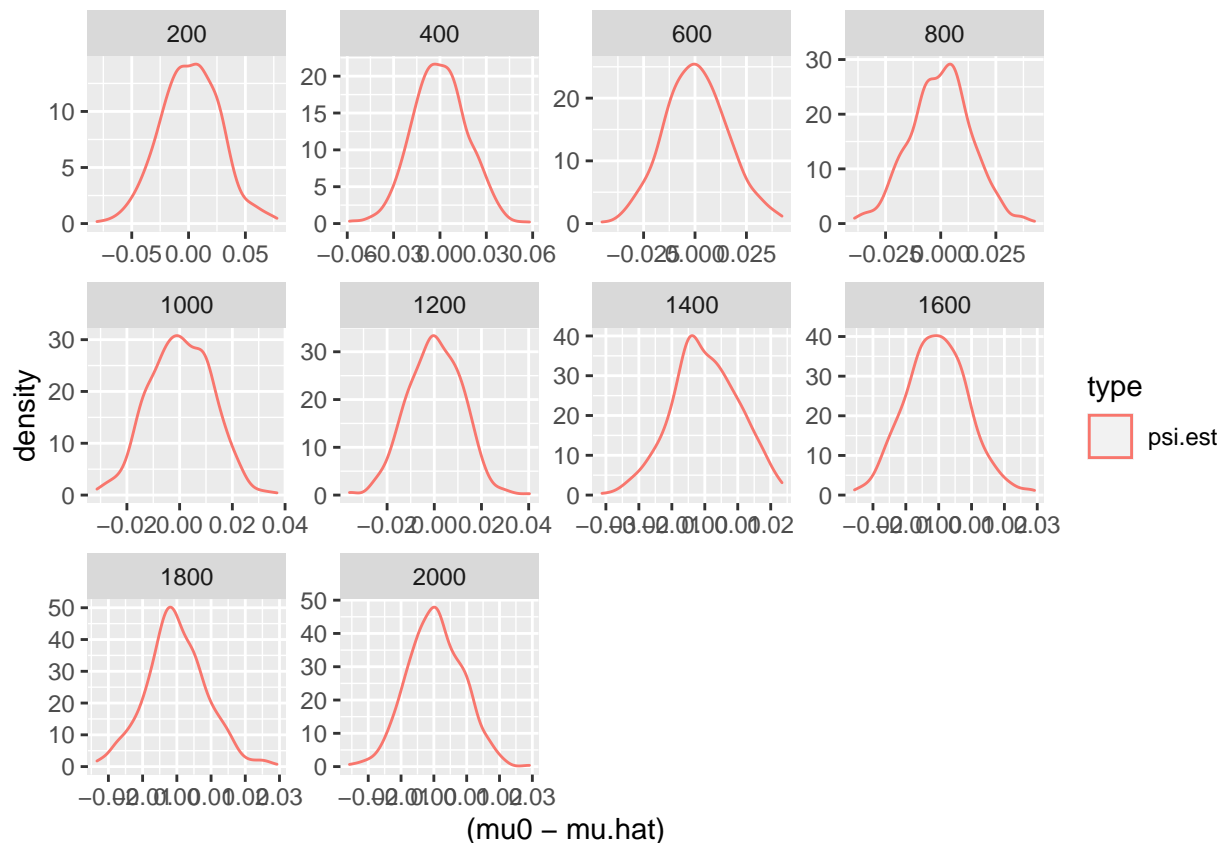
- Both  $\psi_n$  and  $\theta_n$  didn't seem to converge to Normal distribution as  $n$  goes to infinity.

How about the performance of nuisance parameters?

```
ggplot(subset(ests.sim.1, (type %in% 'psi.est'))) +
  geom_density(aes((pi0 - pi.hat), color=type)) +
  facet_wrap(~n, scales='free')
```



```
ggplot(subset(ests.sim.1, (type %in% 'psi.est'))) +
  geom_density(aes((mu0 - mu.hat), color=type)) +
  facet_wrap(~n, scales='free')
```



## 1.2 estimate and confidence interval

```
psi.summaries.1 <- ddply(subset(ests.sim.1, (type %in% 'psi.est')), .(n, type), summarize, na = sum(is.na(coverage)),
  coverage = mean(ll <= psi0 & psi0 <= ul, na.rm=TRUE),
  bias = mean(est - psi0, na.rm=TRUE),
  var = var(est, na.rm=TRUE),
  mse = mean((est - psi0)^2, na.rm=TRUE))
```

### 1.2.1 psi Wald-type

```
library(ggplot2)
p1 <- ggplot(psi.summaries.1) +
  geom_line(aes(n, coverage, color=type)) +
  geom_hline(yintercept=.95)


p2 <- ggplot(psi.summaries.1) +
  geom_line(aes(n, sqrt(n) * bias, color=type))

p3 <- ggplot(psi.summaries.1) +
  geom_line(aes(n, n * mse, color=type))

p4 <- ggplot(psi.summaries.1) +
  geom_line(aes(n, n * var, color=type)) +
  coord_cartesian(ylim=c(0,1))


library(gridExtra)
```

```
grid.arrange(p1, p2, p3, p4, ncol=2)
```



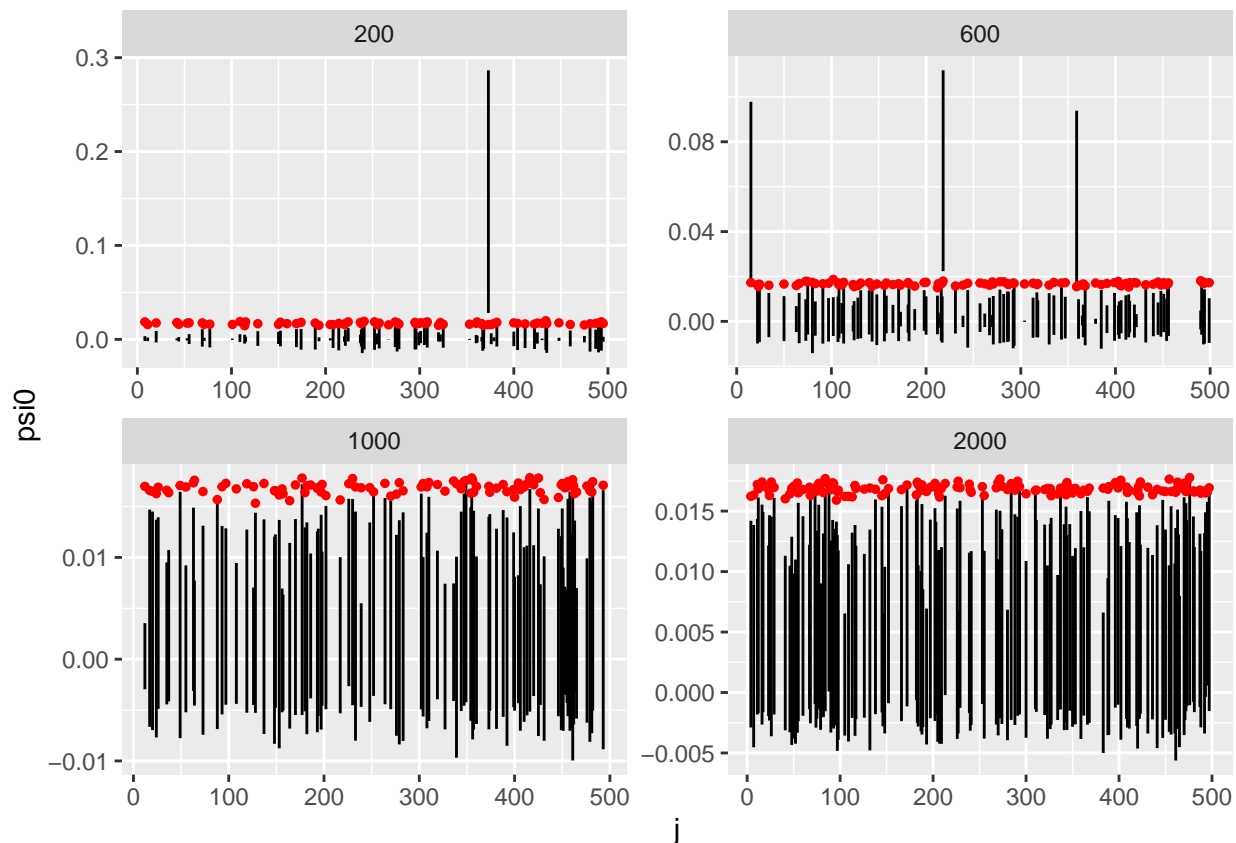
```
ggplot(psi.summaries.1) +  
  geom_line(aes(n, bias, color=type))
```





- The Wald-type CI coverages are below 0.95.
- As  $n$  increases, the coverage is going down.
- (have already check the formula and the formula should be correct).
- The absolute value of bias is getting larger even if plotting with bias instead of  $\sqrt{n} \cdot \text{bias}$ .

```
ests.sim.1$psi.coverage <- (ests.sim.1$ll <= ests.sim.1$psi0) & (ests.sim.1$psi0 <= ests.sim.1$ul)
ests.sim.1$theta.coverage <- (ests.sim.1$ll <= ests.sim.1$theta0) & (ests.sim.1$theta0 <= ests.sim.1$ul)
ggplot(subset(ests.sim.1, (type %in% 'psi.est') & (n %in% c(200, 600, 1000, 2000)) &
  (psi.coverage ==FALSE))) +
  geom_linerange(mapping = aes(x=j, ymin=ll, ymax=ul)) +
  geom_point(mapping=aes(x=j, y=psi0), size=1, col="red") +
  facet_wrap(~n, scales='free')
```




- So I plot out all the Confidence Intervals that fails to cover the  $\psi_0$ .
- The estimator  $\psi_n$  underestimate the real  $\psi_0$ .

### Bootstrap

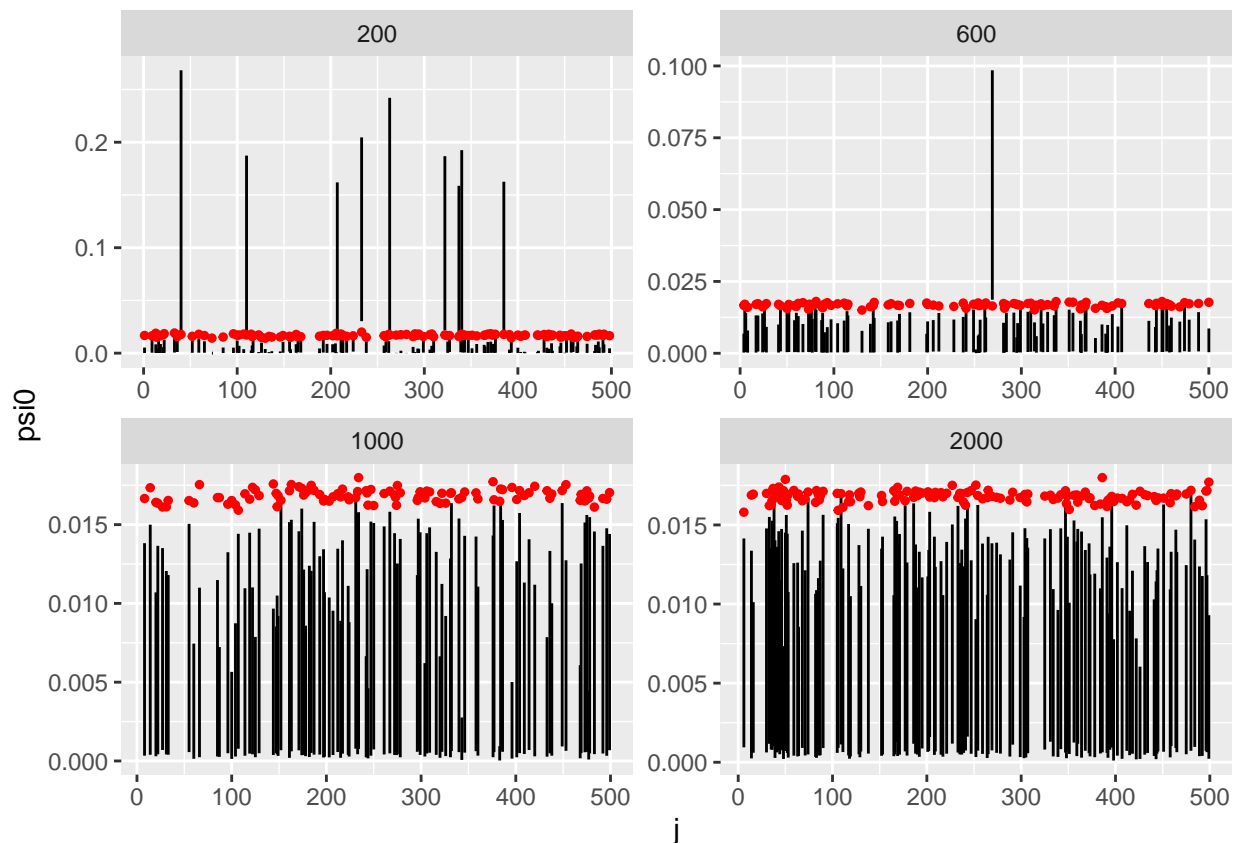
```
# ests.sim.2 <- ests.sim(200*c(1:10), 1:500, control = list(conf.int = TRUE, conf.int.type = 'boot'), n
load("ests.sim.2.Rdata")
# analysis
psi.summaries.2 <- ddply(subset(ests.sim.2, (type %in% 'psi.est')), .(n, type), summarize, na = sum(is.na(
  coverage = mean(ll <= psi0 & psi0 <= ul, na.rm=TRUE),
  bias = mean(est - psi0, na.rm=TRUE),
  var = var(est, na.rm=TRUE),
  mse = mean((est - psi0)^2, na.rm=TRUE))

ggplot() +
  geom_line(data = psi.summaries.1, aes(n, coverage, colour='wald-type')) +
  geom_line(data = psi.summaries.2, aes(n, coverage, colour='bootstrap')) +
  geom_hline(yintercept=.95)
```



- The Bootstrap CI coverages also below 0.95.

```
ests.sim.2$psi.coverage <- (ests.sim.2$ll <= ests.sim.2$psi0) & (ests.sim.2$psi0 <= ests.sim.2$ul)
ests.sim.2$theta.coverage <- (ests.sim.2$ll <= ests.sim.2$theta0) & (ests.sim.2$theta0 <= ests.sim.2$ul)
ggplot(subset(ests.sim.2, (type %in% 'psi.est') & (n %in% c(200, 600, 1000, 2000)) &
  (psi.coverage ==FALSE))) +
  geom_linerange(mapping = aes(x=j, ymin=ll, ymax=ul)) +
  geom_point(mapping=aes(x=j, y=psi0), size=1, col="red") +
  facet_wrap(~n, scales='free')
```



- There are more Bootstrap confidence intervals overestimating the  $\psi_0$  when  $n$  is small.

### 1.2.2 theta Wald-type

```
theta.summaries.1 <- ddply(subset(ests.sim.1, (type %in% 'theta.est')), .(n, type), summarize,
  na = sum(is.na(est)),
  coverage = mean(ll <= theta0 & theta0 <= ul, na.rm=TRUE),
  bias = mean(est - theta0, na.rm=TRUE),
  var = var(est, na.rm=TRUE),
  mse = mean((est - theta0)^2, na.rm=TRUE))





p1 <- ggplot(theta.summaries.1) +
  geom_line(aes(n, coverage, color=type)) +
  geom_hline(yintercept=.95)

p2 <- ggplot(theta.summaries.1) +
  geom_line(aes(n, sqrt(n) * bias, color=type))


p3 <- ggplot(theta.summaries.1) +
  geom_line(aes(n, n * mse, color=type))

p4 <- ggplot(theta.summaries.1) +
  geom_line(aes(n, n * var, color=type)) +
  coord_cartesian(ylim=c(0,1))

library(gridExtra)
grid.arrange(p1, p2, p3, p4, ncol=2)
```




```
ggplot(theta.summaries.1) +
  geom_line(aes(n, bias, color=type))
```



### Bootstrap


```
theta.summaries.2 <- ddply(subset(ests.sim.2, (type %in% 'theta.est')), .(n, type), summarize,
  na = sum(is.na(est)),
  coverage = mean(ll <= theta0 & theta0 <= ul, na.rm=TRUE),
  bias = mean(est - theta0, na.rm=TRUE),
  var = var(est, na.rm=TRUE),
  mse = mean((est - theta0)^2, na.rm=TRUE))

ggplot() +
  geom_line(data = theta.summaries.1, aes(n, coverage, colour='wald-type')) +
  geom_line(data = theta.summaries.2, aes(n, coverage, colour='bootstrap')) +
  geom_hline(yintercept=.95)
```



- Bootstrap confidence levels seem to be weird.

```
ggplot(subset(ests.sim.2, (type %in% 'theta.est') & (n %in% c(200, 600, 1000, 2000)) &
  (psi.coverage == FALSE))) +
  geom_linerange(mapping = aes(x=j, ymin=ll, ymax=ul)) +
  geom_point(mapping=aes(x=j, y=psi0), size=1, col="red") +
  facet_wrap(~n, scales='free')
```



Why Bootstrap confidence levels seem to be so weird?

```
# unit test
null.sims <- FALSE

n <- 1000
set.seed(39338064)
W <- matrix(runif(n*3, 0, 1), ncol=3)
A <- rbinom(n, size = 1, prob = pi0(W))
if(null.sims) {
  Y <- rbinom(n, size = 1, prob = mu0.null(A, W))
  psi0 <- mean((mu0.null(1,W) - mu0.null(0,W))^2)
  theta0 <- var((mu0.null(1,W) - mu0.null(0,W)))
} else {
  Y <- rbinom(n, size = 1, prob = mu0(A, W))
  psi0 <- mean((mu0(1,W) - mu0(0,W))^2)
  theta0 <- var((mu0(1,W) - mu0(0,W)))
}

# use `glm` for both pi.hat and mu.hat in simple unit test
prop.reg <- gam(A ~ s(W[,1]) + s(W[,2]) + s(W[,3]), family = 'binomial')
pi.hat <- prop.reg$fitted.values
AW <- cbind(A, data.frame(W))
mu.reg <- glm(Y ~ ., data=AW, family='binomial')
mu.hat <- mu.reg$fitted.values
mu1.hat <- predict(mu.reg, newdata = cbind(data.frame(A = 1),
                                           data.frame(W)), type = 'response')
```



```

mu0.hat <- predict(mu.reg, newdata = cbind(data.frame(A = 0),
                                           data.frame(W)), type = 'response')
mu.hats <- data.frame(mu1=mu1.hat, mu0=mu0.hat)
# wald
ret1 <- htem.estimator(A, W, Y, control = list(conf.int = TRUE,
                                              pi.hat = pi.hat, mu.hats = mu.hats))

# bootstrap
ret2 <- htem.estimator(A, W, Y, control = list(conf.int = TRUE, conf.int.type='boot',
                                              pi.hat = pi.hat, mu.hats = mu.hats))

cat('theta0:', theta0)

## theta0: 0.003143182
# why bootstrap CI didn't work?
theta.test.vector <- rep(0, 500)
for (i in 1:500){
  boot.inds <- sample(1:n, n, replace=TRUE)

  boot.pi.hat <- pi.hat[boot.inds]
  boot.mu.hats <- mu.hats[boot.inds,]
  boot.ret <- htem.estimator(A[boot.inds], W[boot.inds], Y[boot.inds],
                           control = list(pi.hat = boot.pi.hat,
                                          mu.hats = boot.mu.hats,
                                          conf.int = FALSE))

  theta.test.vector[i] <- boot.ret$est[2]
}
density(theta.test.vector)

##
## Call:
## density.default(x = theta.test.vector)
##
## Data: theta.test.vector (500 obs.); Bandwidth 'bw' = 1.315e-05
##
##          x              y
## Min.   :0.0001208   Min.   : 0.000
## 1st Qu.:0.0008961   1st Qu.: 0.000
## Median :0.0016715   Median : 0.000
## Mean   :0.0016715   Mean    : 322.277
## 3rd Qu.:0.0024468   3rd Qu.: 3.512
## Max.   :0.0032221   Max.    :7908.890

quantile(theta.test.vector, c(0.025,0.975))

##          2.5%          97.5%
## 0.001382813 0.001559773

```

- Since we do bootstrap based on pi.hat and mu.hats, so the estimator won't vary too much.
- The estimator itself also always underestimate the true value.
- Hence it is difficult to cover the true value.

## 2. case 2

- `null.sims=TRUE` it means that `mu0` has nothing to do with A